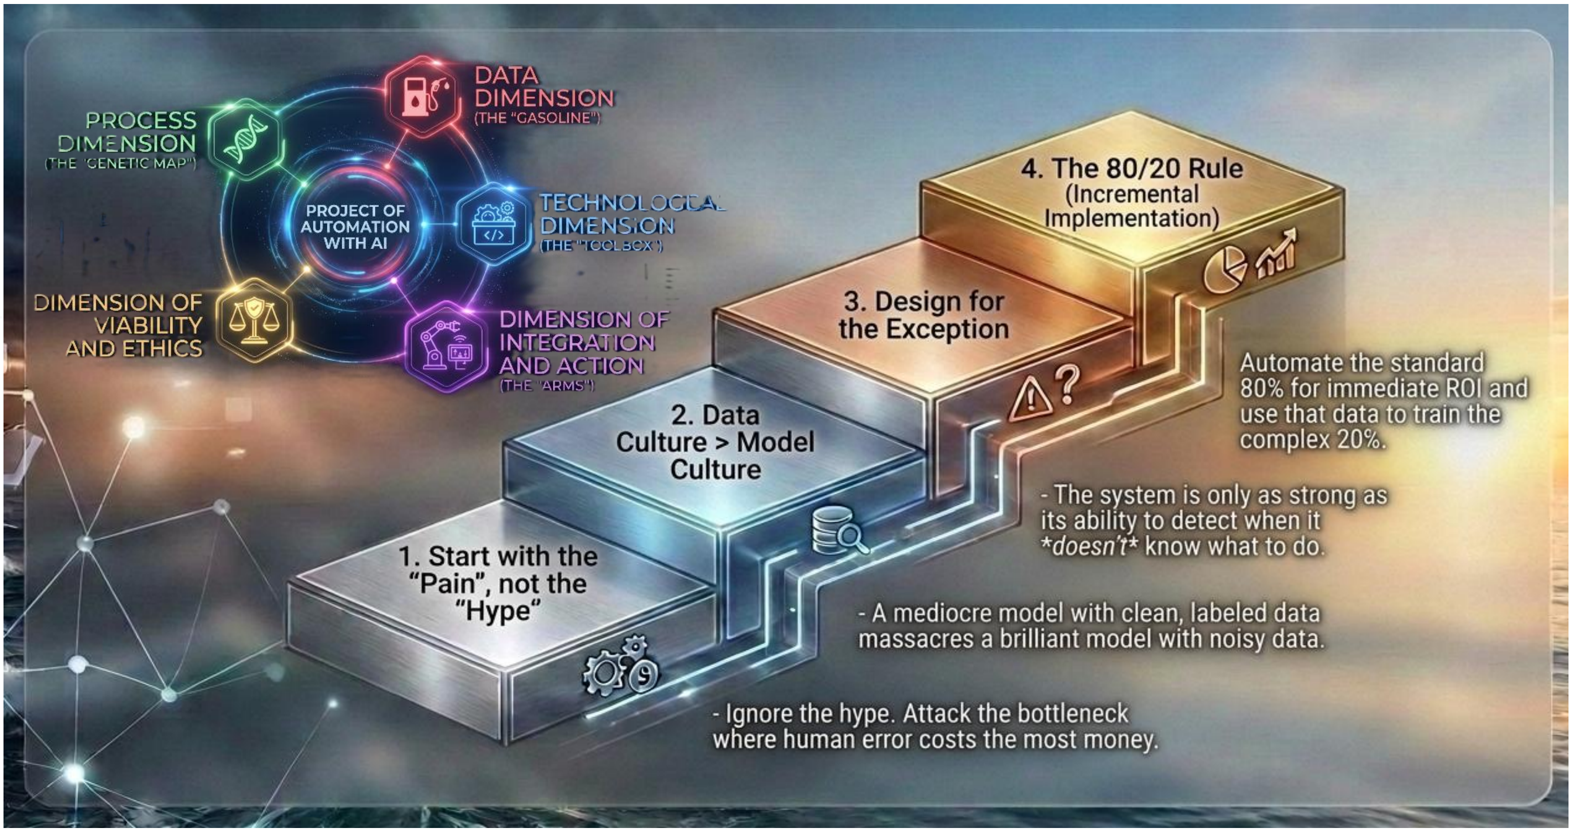

How to Prepare for an AI Automation Project?

"The Five Dimensions of Artificial Intelligence for Automation Model"

Federico Jover Sanz Pastor

April 2026

Originally written in Spanish.

This English version has been translated and reviewed with AI assistance

Content

0. Introduction....4
1. Process Dimension (The “Genetic Map”) 6
1.1. Atomic Decomposition (Input–Process–Output)....6
1.2. Uncertainty Mapping (Edge Cases)....6
1.3. Decision Hierarchy: Who Is in Charge?....7
1.4. “Data Friction” Analysis....7
1.5. The Genetic KPI: Time vs. Quality vs. Cost....7
Practical Tool: The “Process Walkthrough”....8
“Genetic Map” Checklist....8
2. Data Dimension (The “Fuel”)....9
2.1. Typology and Structure (The “Octane”)....9
2.2. The Quality Triad (Accuracy, Completeness and Timeliness)....9
2.3. Availability and Accessibility (The “Logistics”)....10
2.4. Labeling and Preparation (Preprocessing)....10
2.5. Ethics, Privacy and Governance....10
“Fuel” Checklist (Summary)....11
3. Technological Dimension (The “Toolbox”)....12
3.1. Inference Architecture: Where Does the “Brain” Live?....12
3.2. Model Selection: The AI “Zoo”....12
3.3. Orchestration and Agents: The System “Arms”....13
3.4. Connectivity “Middleware” (APIs and Webhooks)....13
3.5. Security and Technological Governance....14
3.6. AI SecOps & Model-Specific Security....15
“Toolbox” Technical Checklist....15
4. Integration and Action Dimension (The “Arms”)....16
4.1. Agency Capabilities (Tool Use / Function Calling)....16
4.2. Systems Orchestration: The “Glue” (Middleware)....16
4.3. Security and Control Protocols (The “Brakes”)....17
4.4. Traceability and Action Logs....17
4.5. Operational Risk and Resilience....18
4.6. Types of Error You Must Anticipate....18

Integration “Arms” Checklist 18
5. Viability and Ethics Dimension 20
5.1. Economic Viability (The ROI of Intelligence) 20
5.2. Algorithmic Ethics and Bias (Fairness) 20
5.3. Governance and Legal Compliance (Compliance) 21
5.4. Risk Transfer: Insurance Coverage and Liability in AI 21
5.5. Human Impact and Change Management 21
5.6. Audit: What to Prepare Before Starting 22
Viability and Ethics Checklist 22
6. Conclusions 23
C1. Final Reflection: The Role of AI as Copilot 23
C2. The “Hybrid Orchestra” Paradigm 23
C3. Roadmap: Global Recommendations 23
C4. Final Launch Checklist (The “Go/No-Go”) 24
C5. What’s Your Next Step? 27
7. Terms of Use and Contact 28
Intellectual Property and CC BY-NC 4.0 License 28
Collaboration and Projects 28
Informal Contact 28

0. Introduction

How to prepare for a complex automation project with AI?

To tackle a **complex automation** project, the leap into the unknown is usually the biggest risk. A "complex" automation is defined as one that has multiple steps, integrates different types of data (vision, text, tables, touch, other physical analysis devices), requires autonomous decision-making and performs physical actions (movement, voice, connection with other physical devices) or logical actions.

You must review the critical dimensions, integrating process analysis and model selection in the early phases of the project so you can decide to approach it in a safer way.

If what you are doing is taking an engine, a base platform with some configuration and you want it to do 4 functions, don't overcomplicate it, but you are in the hands of whoever designed that solution. That is neither good nor bad. It can be efficient; just reflect on the risk you are assuming in your data, in exception handling and in the quality of the deliverable produced by the process. Reading this can help you reflect on these topics.

1. **Process Dimension (the "Genetic Map")**: It consists of dissecting the workflow into logical steps of perception and decision, identifying where AI autonomy ends and human oversight begins. Its success depends, in addition to the precision of the flow, on having exceptions (edge cases) mapped and on defining clear metrics to measure return on investment.
2. **Data Dimension (the "Fuel")**: It focuses on ensuring that the information flow is high quality, properly labeled and representative of business reality, being the main key to success. It is the foundation that prevents AI from hallucinating or inheriting biases, while always guaranteeing legal compliance and privacy.
3. **Technological Dimension (the "Toolbox")**: It is about choosing the "brain" and the right location (Cloud or Edge), balancing power and cost through a mix of heavy and light models. It is the strategic decision on which technical architecture (agent, bot or mathematical model) will best solve each task.
4. **Integration and Action Dimension (the "Arms")**: It defines how AI interacts with your systems (ERP, CRM), or with the outside world (machinery, devices, robots, voice, etc.) and what permissions it has to execute real changes. It must necessarily include a manual fallback plan to act immediately if the technology fails or the model makes a mistake, ensuring business resilience.
5. **Viability and Ethics Dimension**: It assesses whether the project is economically sustainable in the long term in terms of compute cost, and whether its decisions are explainable and fair. It seeks to prevent automation from becoming an uncontrollable "black box" or an unsustainable operational expense.

Unfortunately, problems **will** arise in projects, and the real merit lies in solving them quickly without losing focus on the objectives by getting stuck in technical details.

Before talking about technology, assume one practical rule: the "minimum product" of an AI automation project is not the model, it is an end-to-end functional flow, even if it only covers a subset of the process (20–30% of the real volume). Your initial goal is not to automate everything, but to prove that AI can execute the happy path with acceptable quality, traceability and safety.

Priority Summary

Priority	Element	Why is it critical?
Critical	Clean Data	Without this, even the most advanced AI will produce garbage results (*Garbage In, Garbage Out*).
High	Exception Map	This is where automation projects fail and costs skyrocket.
Medium	Model Selection	It affects cost and speed, but it is easier to iterate.

PRO Tip: Before starting development, try to run the "automated" process manually, but strictly following the logic you would give to the AI. If a human gets confused with your instructions, the AI will too.

And remember:

- We are going to look at each of these dimensions in more detail so that it helps you define a project where the problems that arise later are manageable.
- But the project plan, with its waterfall Gantt, Agile methodology or whatever you choose, still has to be done. These dimensions are a support for what you must consider in your plans.
- **In the full project plan you must integrate detailed analysis, development, deployment and CHANGE MANAGEMENT.**

1. Process Dimension (The “Genetic Map”)

Diving deep into the **Process Dimension** is likely the factor that determines whether an AI project saves millions or becomes a bottomless pit of engineering hours. In 2026, we no longer design “workflows”; we design **“decision ecosystems”**.

To map the “DNA” of your process, you must break it down into several layers of technical–operational analysis:

1.1. Atomic Decomposition (Input–Process–Output)

Do not look at the process as a straight line; see it as a series of chemical transformations. Each step must be defined by the **mutation of the information**:

- **Input Analysis:** Does the data arrive structured (JSON, CSV) or unstructured (a WhatsApp voice note from an angry customer, a photo of a crumpled delivery note)?
  - *If it is unstructured, you need a prior **Perception** layer (OCR, NLP).*
- **Transformation Analysis:** What happens to the data?
  - **Extraction:** Pull out names and dates.
  - **Validation:** Check whether the data is logical.
  - **Inference:** Predict whether the customer is likely to churn based on that data.
- **Output Analysis:** What does the step produce? A new data point for the next step or a final action (send an email)?

In processes involving physical interaction, the map must distinguish between deterministic events (sensor reading) and probabilistic events (image interpretation by AI). It is vital to map not only software failure, but hardware degradation: what does the process do if a camera goes out of focus or a sound sensor picks up excessive background noise?

1.2. Uncertainty Mapping (Edge Cases)

The biggest enemy of AI is not error, it is **ambiguity**. You must create an inventory of exceptions:

- **Happy Path:** The 70–80% of cases where everything is standard. This is where AI flies.
- **Edge Cases:** Physical and Logical Exceptions Matrix:
  Do I have 90% of out-of-normal cases documented, including sensor failures (blurry cameras, background noise) and physical obstructions? What does the system do if the customer sends a photo of a product that is not ours? And if the text is in a language the model does not handle well?
- **Escalation Protocol:** You must define the **“Confidence Threshold”**.
  - *Example:* If the model is less than 85% confident that a claim is valid, it **does not automate**; it routes the case to a human.

To ground the concept, here are real examples of how Edge Cases appear in different industries:

- **Logistics:** Photos of parcels taken in poor lighting, partially torn labels, incomplete addresses, or customers attaching images of products that are not from the company.
- **Manufacturing:** Miscalibrated vibration sensors, dusty cameras, parts partially covered by operators, or color variations due to lighting changes.

- **Finance:** Scanned documents with shadows, illegible signatures, numeric fields with inconsistent formats (“1.200”, “1,200”, “1200”).
- **Customer Service:** Audio with background noise, messages in mixed languages, irony or sarcasm that the model may interpret literally.

These examples should be incorporated into the Exceptions Matrix to avoid the system failing in situations which, although rare, happen often enough to break an automation.

Beyond listing them, you should quantify how often each of these edge cases appears in your real operation and what impact they have when they are mishandled. A single rare but critical pattern (for example, a specific combination of bad lighting and damaged labels in logistics) can invalidate the whole automation if it is not explicitly modelled. Treat these edge cases as first-class citizens in your design, not as footnotes: they are precisely where projects usually break in production.

1.3. Decision Hierarchy: Who Is in Charge?

In complex automation, you must decide the level of autonomy for each node in the genetic map:

1. **Decision Support (Human-in-the-loop):** The AI analyzes and proposes 3 options; the human chooses one. Ideal for legal or medical processes.
2. **Active Supervision (Human-on-the-loop):** The AI executes the action, but a human can cancel it within a “grace” period (e.g. 5 minutes before sending a bank transfer).
3. **Full Autonomy (Human-out-of-the-loop):** For high-volume, low-risk processes (e.g. spam classification or real-time micro price adjustments).

1.4. “Data Friction” Analysis

This point is crucial: how many times does the data have to “jump” from one system to another?

- Each jump is a potential failure point for an **AI Agent**.
- **Latency Check:** If the process requires a vision model to analyze real-time video in a factory, you cannot afford for the process map to depend on a cloud API that takes 2 seconds to respond. The genetic map will tell you that this step must be **Edge AI** (local processing).

1.5. The Genetic KPI: Time vs. Quality vs. Cost

For the analysis to be complete, each step of the process must have its current cost assigned:

- **Error Cost:** How much does it cost us if a human makes a mistake at this step? Sometimes the AI is “worse” than the best human, but “better” than the average, and that is already profitable.
- **Inertia Cost:** How much do we lose for every hour this process is stopped waiting for a signature?

Practical Tool: The "Process Walkthrough"

Before you write any code, do some **shadowing**: sit next to the person who does the work today.

- Do not ask them what they do (they will tell you the manual).
- Watch what they do when something goes wrong. Those "hacks" and manual workarounds are exactly what your AI will need to learn to handle – or what will make your project fail if you do not map them.

"Genetic Map" Checklist

1. Have I broken the process down into clear Input → Transformation → Output steps, with no "magic" jumps?
2. Do I have the "Happy Path" (70–80% of standard cases) identified as the first target for automation?
3. Have I listed the main edge cases and defined what should happen to them (route to human, ask for more data, block)? Does this include the Physical and Logical Exceptions Matrix, including sensor failures (blurry cameras, background noise) and physical obstructions?
4. Have I defined the level of autonomy for each step: decision support, active supervision, or full autonomy?
5. Have I mapped the data friction points (system hops, critical latencies, problematic formats)?
6. Do I know the error cost and inertia cost of each key step in the process?
7. Have I done at least one "Process Walkthrough" (shadowing), observing what people do when something goes wrong?
8. Can the process, as described, be followed from beginning to end without someone saying "it depends" every two steps?

2. Data Dimension (The “Fuel”)

The **Data Dimension** is where most AI projects die. It is not just about “having data”, it is about having it in a form that AI can digest and that does not mislead it. In complex automation, data is not static; it flows, changes format and loses quality over time.

The quality of sensor data (images, audio, vibration) critically depends on environmental conditions. You must set signal-cleanliness standards: constant lighting for computer vision and frequency isolation for sound interpretation. Without normalized physical data, the model will generate constant false positives.

To break down this dimension, we need to analyze five critical axes:

2.1. Typology and Structure (The “Octane”)

Not all data is processed in the same way. You must classify your “fuel” to know which “engine” (model) to use:

- **Structured Data (Tabular):** SQL databases, Excel, financial logs.
  These are ideal for **classical Machine Learning** (XGBoost, Decision Trees). They are predictable and easy to validate.
- **Unstructured Data:** Free text (emails), audio, images or video.
  This is where **LLMs and Deep Neural Networks** come into play. The challenge here is “cleaning” (removing background noise from audio or ads from text).
- **Semi-structured Data:** JSON, XML, HTML.
  They require a **parsing** layer (syntactic analysis) before the AI can reason about them.

2.2. The Quality Triad (Accuracy, Completeness and Timeliness)

A model with bad data is a dangerous model. You must audit:

- **Accuracy (Ground Truth):** Who labeled this data?
  If you use historical human data where people made mistakes, the AI will learn and automate those mistakes.
  You need a **“Ground Truth”** set to validate against, and you must also ask: have we calibrated the sensors? Have I validated that training data matches real conditions (light, noise, temperature) and that the hardware is calibrated?
- **Completeness (Missing Values):** Predictive models hate gaps.
  If 40% of customers in your database are missing “age”, the model will bias its predictions. You must decide whether to impute (fill with means) or discard.
- **Timeliness (Data Drift):** Data expires.
  A purchase-behavior model from 2019 is not fit for 2026.
  You must ensure that AI is trained on data reflecting the current reality. In practice, this means you also have to define the process to periodically check whether the patterns the model learned are still valid and, if not, retrain it with more recent data.

To these three classical dimensions you can add a fourth: **cross-source consistency**. If CRM, ERP and Excel say different things about the same customer or order, the AI must not decide on its own which one is the truth without an explicit business rule.

Include master-source detection and priority rules between systems. When two systems disagree, there must be a formal hierarchy determining which one prevails.

2.3. Availability and Accessibility (The "Logistics")

This is the most common technical failure: the data exists, but AI cannot reach it.

- **Data Silos:** Is the data in a CRM that AI cannot access for security reasons?
- **Access Latency:** If the automation is real-time, can you extract the data in milliseconds, or does the database query take 10 seconds?
- **APIs and Connectors:** Are there technical "bridges" for the model to read the data automatically, or do you depend on a manual CSV export every Monday?

2.4. Labeling and Preparation (Preprocessing)

For complex automations (such as computer vision or contract classification), the data needs "instructions":

- **Data Labeling:**
  If you want AI to detect cracks in pipes, you need thousands of photos where an expert has drawn a circle around the crack.
  This is the most expensive and slowest part of the project.
- **Normalization:**
  Ensuring that "100€", "100 euros" and "100.00 EUR" are processed as the same numeric value by the model.
- **Data Augmentation:**
  If you have few examples, you can use techniques to create artificial variations (e.g. rotate images, change synonyms in text) so the model becomes more robust.

2.5. Ethics, Privacy and Governance

By 2026, legal compliance is part of data engineering:

- **Anonymization / Pseudonymization:**
  Can the model learn without seeing real names or ID numbers? (Use of **PII Masking**).
- **Selection Bias:**
  If you are going to automate hiring and your historical data only has men in senior roles, AI will reject women by "statistical pattern". You must rebalance the dataset.
- **Data Lineage:**
  You must be able to trace where a piece of data came from if AI makes a wrong decision that causes a legal issue.

“Fuel” Checklist (Summary)

1. Do I have a single source of truth for each critical data point?
2. Have I identified the missing data and do I have a plan for it?
3. Does the data arrival speed match the speed that AI needs?
4. Is the data protected according to privacy regulations?
5. **Sensor Quality and Calibration:** Have I validated that training data matches real-world conditions (light, noise, temperature) and that the hardware is calibrated?
6. **Fire test:** Could a human expert perform the task using *exactly* the same data I am going to give to the AI? If the answer is NO, the AI will not be able to either.

3. Technological Dimension (The "Toolbox")

The **Technological Dimension** is the bridge between process design and operational reality. In complex automation, you do not choose a single tool; you build a **stack** of technologies that must communicate with each other without friction.

In 2026, this dimension is divided into four critical layers you should define before writing a single line of code:

3.1. Inference Architecture: Where Does the "Brain" Live?

The choice between Cloud, On-premise or Hybrid defines the speed, cost and security of your project.

- **Cloud AI:** Using models via API (such as Claude on AWS Bedrock or GPT on Azure).
  - *Advantage:* Access to the most powerful models in the world without investing in hardware.
  - *Risk:* Dependency on the internet and variable usage-based costs (tokens).
- **Edge AI (On-device/Local):** The model runs on a chip inside a camera, a sensor or a server within the company.
  - *Advantage:* Zero latency (instant response) and full privacy.
  - *Use cases:* Detecting failures on a production line or autonomous drones.
- **Hybrid:** A small local model (SLM) filters data and only sends the complex part to the cloud.

For physical actions:

- Industrial protocols must be followed: physical connectivity standards such as MQTT or OPC-UA for communication with PLCs and heavy machinery.
- Edge AI architecture is not optional; it is mandatory. Cloud latency is usually unacceptable for stopping a robotic arm or a conveyor belt. Specialized models (CNNs, ViTs) must run on local hardware to guarantee millisecond response.

3.2. Model Selection: The AI "Zoo"

Do not use a cannon to kill a fly. You must assign the right model to each step in your genetic map:

- **Frontier Models:** Models like Claude 4 Opus or GPT-5. Used for **complex reasoning**, strategic planning or when error is unacceptable.
- **Small Models (SLMs):** Models like Llama 3 (8B) or Phi-4. Ideal for specific tasks (summarization, classification, data extraction) because they are extremely fast and cheap.
- Specialized Models (Non-LLMs):
  - **Vision Transformers (ViT):** When your automation needs to "see" (part quality, safety).
  - **Graph Models:** When your problem involves complex relationships (fraud networks, logistics networks).

A simple rule: always start with the smallest, cheapest model that can solve the problem. Only move up to frontier models when the use case justifies it (high risk, extreme complexity or very strict quality requirements).

3.3. Orchestration and Agents: The System "Arms"

This is where AI stops "talking" and starts "doing". In complex automation, you need tools that manage the flow:

- **Agent Frameworks (e.g. LangChain, CrewAI, AutoGen):** They allow several models to collaborate. A "researcher" agent fetches data, and a "writer" agent produces the report.
- **Memory Systems (Vector Databases):** Tools like Pinecone or Weaviate. They allow AI to "remember" company documents or past interactions so it does not hallucinate and always uses real information (**RAG – Retrieval-Augmented Generation**).
- **Execution Tools (Tools/Functions):** The model's ability to call a code function, run a Python script or query an SQL database autonomously.

Integration must be idempotent: every command sent to a physical device must be safely repeatable without causing damage (e.g. not trying to close a valve that is already closed). A handshaking protocol is needed between AI and hardware (PLC/IoT) to confirm that the physical action has completed before proceeding to the next step.

Not every automation needs an agent.

- Use a **model** when the task is to classify, extract, predict or generate.
- Use a **bot** when the interaction is repetitive and low-reasoning.
- Use an **agent** when there are multiple steps, external tools and intermediate decisions.
- Use **RPA** when there is no API and you must operate through legacy interfaces.

A useful mental shortcut is to **start from the simplest option** and only move up the ladder when there is a clear reason. If a single model with a couple of functions can solve the use case, you probably do not need an agent at all. If you are **tempted to introduce several agents talking to each other, stop and ask: "What concrete failure do they solve that a simpler orchestration cannot?"** In many real projects, the most robust solutions come from keeping the number of moving parts as low as possible and reserving multi-agent setups only for genuinely open-ended or multi-tool problems.

3.4. Connectivity "Middleware" (APIs and Webhooks)

Complex automation usually fails not because of AI, but because of the "pipes" that connect legacy systems with modern AI.

- **RPA (Robotic Process Automation) with AI:** If your company uses old software without APIs, you need a "robot" to move the mouse on behalf of the AI.
- **Event-Driven Architecture:** The system should not be polling "is there anything new?" every second. It should work by events: "When an email arrives (Event) → Trigger the AI Agent".

- **API Management:** You need a layer that controls how many requests you make to models so you do not overload the system or blow up the bill.

You also need to define **integration patterns**, because they determine whether your AI automation will be **stable, scalable and safe**, or whether it will become a fragile system that fails precisely when you need it most.

In complex automation, AI does not work alone: it depends on ERPs, CRMs, databases, sensors, APIs, message queues, webhooks… and the way all these elements are connected often **defines the success of the project more than the model itself**.

Typical patterns:

- **Event-Driven:** Ideal for reactive processes (incoming payments, incoming emails).
- **Centralized Orchestration:** One master system coordinates every step (useful for complex processes with dependencies).
- **Choreography:** Each system reacts to events without a central controller (more scalable but harder to debug).
- **Circuit Breakers:** Prevent a failure in an external system from collapsing the whole automation.
- **Asynchronous Queues:** Allow you to absorb load spikes without losing data.

These patterns drastically reduce system fragility and prevent bottlenecks.

3.5. Security and Technological Governance

- **Guardrails:** Software layers that review what AI is about to say or do before it happens, blocking sensitive content or unauthorized actions.
- **MLOps (ML Operations):** A system to monitor that the model does not lose accuracy over time (*model drift*) and to update it without stopping the business process.
  You must also consider physical sensor wear and tear.
- **Vendor Lock-in:** Assess what happens if the AI provider breaches SLAs, changes prices, limits capabilities or discontinues the model.
  Whenever possible, design an architecture that allows the model to be swapped with minimal changes to the rest of the system.
- **Audit Frequency:** Establish criteria to determine when a model needs retraining based on real performance metrics.
- **Model and Data Versioning:** Emphasize the importance of keeping a historical record (lineage) not only of data, but also of model versions used for each decision.

For an AI system to be sustainable, at minimum you should monitor these operational metrics:

- **Data Drift:** Changes in the input data distribution relative to training.
- **Concept Drift:** Changes in the relationship between data and label (e.g. customer behavior shifts).
- **Per-model Latency:** Actual inference time vs. process SLA.
- **Retry Rate:** Indicates connectivity or stability issues in the system.
- **Fallback-to-Human Rate:** If this goes up, the model is losing accuracy or confidence.
- **Cost per Inference:** Essential to avoid the project becoming economically unviable.

These metrics should appear in operational dashboards and be reviewed periodically.

3.6. AI SecOps & Model-Specific Security

This section focuses on the technical shielding of AI-specific logic. It is critical because standard cybersecurity measures often fail to detect or prevent attacks like prompt injection or data poisoning, which can manipulate the model's behavior directly without breaching the traditional IT perimeter.

- **Defense-in-Depth Architecture:** Implement specialized security layers designed to protect the model from adversarial attacks.
- **Vulnerability Mitigation:** Actively defend against *Prompt Injection* (both direct and indirect), *Jailbreaking*, and *Data Poisoning* that could corrupt the model's logic.
- **Data Privacy & Sanitization:** Deploy automated protocols for PII (Personally Identifiable Information) masking and redaction to prevent data leakage.
- **Guardrail Systems:** Use secondary "Guardrail" models to filter both inputs and outputs, ensuring content stays within safe and ethical bounds.
- **Continuous Monitoring:** Establish a SecOps lifecycle to detect "adversarial drift," where attackers adapt their methods to bypass static filters over time.

“Toolbox” Technical Checklist

1. **Right Model?** Have I justified using an expensive LLM versus a cheaper statistical model for each step?
2. **Latency?** Is the chosen technology’s response time lower than the process’s maximum acceptable time?
3. **Scalability?** If data volume increases 10x tomorrow, will the tech stack cope or collapse?
4. **Security?** Do sensitive data stay inside the company perimeter, or do they leave to public clouds?
5. **Interoperability?** Do all components (AI, database, CRM) speak the same language (REST APIs, JSON)?

Key takeaway: In the Technological Dimension, **“less is more”**.
A robust automation built on a simple model and a good database is preferable to a swarm of complex agents that no one knows how to fix when they fail.

4. Integration and Action Dimension (The "Arms")

The **Integration and Action Dimension** is the moment of truth: it is where AI "thinking" becomes a concrete "action" in the real or digital world. Without this dimension, your AI is just a consultant that gives advice; with it, it is an operator that does the work.

In a complex automation, the "arms" must be precise, safe and perfectly coordinated with existing systems. Here we detail the pillars of this dimension:

4.1. Agency Capabilities (Tool Use / Function Calling)

For AI to act, it needs an interface that allows it to "touch" other programs. This is achieved through **Tool Use**.

- **Tool Definition:** You must explicitly register what the AI is allowed to do (e.g. `send_email`, `check_stock`, `generate_invoice`).
  Each tool is a code function that the model decides to invoke.
- **Parameter Handling:** The model must be able to extract from the conversation the data required by the tool (e.g. extract the "product ID" and the "quantity" to place an order).
- **Feedback Loop:** AI must "read" the result of its action.
  If it tries to create an invoice and the system returns an error "Customer not found", the AI must process that error and decide the next step (e.g. ask the customer for their tax ID).

4.2. Systems Orchestration: The "Glue" (Middleware)

AI rarely connects directly to an ERP. You need an intermediate layer to manage communication.

- **APIs (Application Programming Interfaces):** This is the gold standard. If your systems (SAP, Salesforce, Zendesk) have modern APIs, integration is smooth.
- **Webhooks:** They allow your systems to "notify" AI when something important happens (e.g. "A payment was received") without the AI constantly polling.
- **RPA (Robotic Process Automation):** For legacy systems that have no API. Here, AI controls software that mimics mouse clicks and keystrokes to interact with old interfaces.

You must also define **integration patterns**, because they determine whether your AI automation will be **stable, scalable and safe**, or whether it will become a fragile system that fails when you need it most.

In complex automation, AI does not work alone: it depends on ERPs, CRMs, databases, sensors, APIs, message queues, webhooks… and the way all these elements are connected often **defines the success of the project more than the model itself**.

Common patterns include:

- **Event-Driven:** Ideal for reactive processes (received payments, incoming emails).
- **Centralized Orchestration:** A master system coordinates every step (useful for complex processes with dependencies).

- **Choreography:** Each system reacts to events without a central controller (more scalable, but harder to debug).
- **Circuit Breakers:** Prevent a failure in an external system from collapsing the entire automation.
- **Asynchronous Queues:** Allow you to absorb spikes in load without losing data.

These patterns dramatically reduce system fragility and avoid bottlenecks.

4.3. Security and Control Protocols (The “Brakes”)

Giving AI “arms” comes with risks. You need mechanisms that prevent disasters:

- **RBAC (Role-Based Access Control):** AI must not have “superpowers”. Its system credentials must be strictly limited to what it needs to do its job (Principle of Least Privilege).
- **Human-in-the-loop (HITL) for Critical Actions:** You must define “red lines”.
  - *Example:* AI can draft the claim email, but **a human must press “Send”**.
    Or AI can approve refunds up to €50, but if they are €500, manual authorization is required.
- **Sandboxing:** Executing code actions in isolated environments so that an AI error cannot wipe a database or compromise the server.
- **Action Idempotency:** AI must not execute the same action twice due to a network failure or retry.
- **Pre-action Validation:** Before acting, the system must check that the context and parameters are still valid.
- **Explicit Confirmation:** For critical actions, AI must wait for technical and functional confirmation before marking the task as completed.
- **Physical Staging Environments (Testing):** You must define how to build test environments that faithfully simulate the physical world or critical systems before going live.

4.4. Traceability and Action Logs

In complex automation, when something goes wrong, you need to know exactly what happened.

- **Audit Trail:** An immutable record of every decision and action:
  “At 10:01 the model decided the customer was VIP based on data X, and at 10:02 it executed the 20% discount function.”
- **Error Handling and Retries:** If a server is down, AI must have an “exponential backoff” retry logic or, ultimately, an escape route to alert human support.
- **State Confirmation:** Never assume an action completed successfully.
  The “arm” must wait for a “200 OK” (or equivalent) from the target system before marking the task as finished.
- **Event Correlation:** Each decision and action must share a unique identifier to reconstruct the complete sequence.
- **Rejected Decisions Log:** Do not only store what AI did, but also what it proposed and was not executed.
- **Retry Traceability:** Log how many times an action was attempted and why it failed.

4.5. Operational Risk and Resilience

Every complex automation project must consider what happens when something fails **outside** the model: API outages, timeouts, network errors, queue saturation, legacy system downtime or degraded performance from an external provider.

This analysis must include, at minimum:

- Definition of expected failure modes.
- Strategy for manual fallback or automated alternative (BCP).
- Retries, queues and circuit breakers.
- Controlled degradation procedures.
- Critical third-party dependencies and their impact on the service.

4.6. Types of Error You Must Anticipate

In practice, almost all problems in a complex automation project can be grouped into four types of error:

- **Perception Errors:** AI “sees” the world incorrectly (wrong OCR, incomplete data extraction).
  These are fixed by tuning the perception layer or asking for additional confirmation.
- **Decision Errors:** AI makes the wrong decision based on correct data (bad model or prompt, poorly calibrated confidence threshold).
  These are mitigated with thresholds, human review and reevaluating the model.
- **Action Errors:** AI executes something it should not, or executes it twice (permission issues, lack of idempotency, missing pre-validation).
  These are controlled with brakes, RBAC and integration testing.
- **Infrastructure Errors:** The system or an external provider fails (APIs down, timeouts, saturated queues).
  These are handled with retries, circuit breakers and controlled degradation plans.

Explicitly reviewing these four error types while designing your solution forces you to translate abstract worries (“what if it fails?”) into concrete conditions and protection mechanisms.

Integration “Arms” Checklist

1. **Tool Catalogue:** Are all functions the AI can use documented and constrained?
2. **Identity Management:** Does AI have its own user and password in corporate systems to track its actions?
3. **Autonomy Limits:** Is there a defined monetary or sensitivity threshold above which AI *must* ask for permission?
4. **Notification System:** How does the human team find out if one of AI’s “arms” gets stuck?
5. **Staging Environment:** Do I have a mirror of my real systems where AI can “break things” while we learn how to integrate it?
6. **Operational Risk and Resilience:** Have I identified the most likely end-to-end failure points? Is there a continuity plan if an AI provider or external API goes down? Can the system degrade without fully stopping operations? Are circuit breakers, queues or retries defined? Is the return to manual mode documented?

7. **Safety and Mechanical Fallback Protocol:** Is there a physical emergency stop independent of AI and a “failsafe” state (safe machinery lock) if control is lost?

Conclusion: Integration is what separates an AI experiment from a **business solution**. If the arms are not well articulated, the brightest brain (the model) will be useless.

5. Viability and Ethics Dimension

The **Viability and Ethics Dimension** is the final filter that determines whether an AI project should exist at all. In 2026, it is no longer enough for something to be technically possible; it must also be economically sustainable, legally safe and socially responsible. A "successful" automation that destroys the company's reputation or burns through the budget in a month is, in reality, a failure.

Viability must include the **TCO of the hardware**: the cost of maintaining cameras, sensors and their periodic cleaning. At an ethical level, explainability must be forensic: in the event of a physical incident, we must be able to reconstruct exactly which pixels or audio frequencies led the AI to activate a device.

Here we detail the four pillars that ensure your project is solid in the long term:

5.1. Economic Viability (The ROI of Intelligence)

Unlike traditional software, AI has variable costs that can scale exponentially. You must calculate the **Total Cost of Ownership (TCO)**:

- **Inference Cost:**
  How much does each execution cost?
  A model like Claude 4 processing thousands of emails can be more expensive than the salary of the person who used to read them.
  You must balance the use of "Premium" models vs. "Light" models.
- **Maintenance Cost (MLOps):**
  Models degrade over time. You need budget to periodically retrain AI with new data.
  If there are physical devices, include **Periodic Cleaning and Calibration**.
- **Opportunity and Error Cost:**
  How much money do we lose if AI makes a serious mistake?
  This "insurance" must be included in the viability analysis.
  For example, in a complaints process, a single error that denies a legitimate refund may cost you less directly than the reputational damage if it goes viral.

5.2. Algorithmic Ethics and Bias (Fairness)

AI is not neutral; it is a mirror of the data it was trained on.
If your historical data contains bias, the automation will execute that bias at industrial scale.

- **Bias Identification:** You must audit whether the model discriminates by gender, age, ethnicity or postal code.
  This is especially critical in HR processes, credit or insurance decisions.
- **Fairness:** Implement stress tests where you change a sensitive variable (e.g. gender in a résumé) and verify that AI's outcome is the same.
- **Transparency:** Avoid the "black box".
  If AI makes a decision that affects a human, you must be able to explain which factors weighed the most in that decision.

5.3. Governance and Legal Compliance (Compliance)

In 2026, the EU AI Act and similar regulations worldwide impose strict rules depending on the risk level of the system.

- **Risk Classification:** Determine whether your project is “High Risk” (healthcare, critical infrastructure, employment, etc.).
  If it is, documentation and audit requirements are massive.
- **Data Sovereignty:** Where does the information travel?
  If you use a US cloud to process data from European citizens, you may violate data protection laws.
- **Intellectual Property:** Ensure that AI-generated content or data used for training does not infringe third-party copyrights.

5.4. Risk Transfer: Insurance Coverage and Liability in AI

This module defines the legal and financial safety nets required for autonomous systems. It is essential because AI 'hallucinations' or errors can lead to real-world damages; having a risk transfer strategy ensures the organization is not financially paralyzed by the inherent unpredictability of the 'black box'.

- **Accountability Framework:** Establish a clear legal and financial boundary for responsibility when AI autonomy leads to errors, hallucinations, or damages.
- **Contractual Liability:** Define specific terms between providers, integrators, and end-users to manage the "Black Box" problem and liability gaps.
- **Specialized AI Insurance:** Evaluate and procure "AI-specific" insurance or cyber-policies that cover algorithmic failures and professional indemnity in automated environments.
- **Financial Mitigation:** Ensure that the potential fallout from incorrect automated decisions is backed by robust risk-transfer strategies and financial reserves.

5.5. Human Impact and Change Management

Complex automation often generates fear and resistance among staff.
A key ethical dimension is how the machine is integrated with people.

- **Upskilling (Reskilling):** What will people whose tasks have been automated do?
  The project is ethically viable if it **transforms roles instead of simply eliminating them**.
- **Accountability:** There must be a human owner (Chief AI Officer or business leader) responsible for each AI decision.
  “The algorithm did it” is no longer a valid defense before a judge or a customer.
- **Energy Sustainability:** Large models consume significant energy.
  Evaluating the environmental impact of your architecture (carbon footprint) is part of modern corporate ethics.

- **Knowledge Concentration Risk:** If today the process depends on one or two key people, the project must explicitly include capturing and documenting their tacit knowledge (unwritten rules, tricks, exceptions).
  Without this transfer, automation will only partially replicate reality and will be fragile in the face of staff turnover.
- **Resistance Management:** Identify and include methods to mitigate fear of being replaced by AI, which is a critical factor in the operational success of these projects.

5.6. Audit: What to Prepare Before Starting

Audit readiness is the proactive organization of all project evidence and logic. It is important because it transforms the AI from a 'mysterious experiment' into a transparent business process, ensuring compliance with global regulations (like the EU AI Act) and facilitating institutional trust.

- **Pre-Audit Repository:** Consolidate all project artifacts (maps, data logs, risk matrices) into a centralized folder to ensure transparency from day one.
- **Key Documentation:** Maintain updated versions of the *Genetic Map* (process flow), *Data Lineage* (provenance and cleaning logs), and the *Risk Matrix*.
- **Decision Log (Audit Trail):** Implement a technical log that records the rationale and confidence scores behind AI-driven outputs for future forensic review.
- **Regulatory Compliance:** Align documentation with international standards (such as the *EU AI Act* or *NIST AI RMF*) to prove the system is predictable, fair, and auditable before it reaches production.

Viability and Ethics Checklist

1. **Cost per Task:** Have I calculated the cost in tokens/compute for a single successful execution?
2. **Bias Audit:** Have I tested the model with diverse data profiles to ensure it does not discriminate?
3. **Panic Button:** Is there a legal and technical procedure to shut down AI if it starts behaving erratically?
4. **Internal Communication:** Do employees know how AI will help them and what will happen to their current tasks?
5. **Compliance:** Do I have legal’s approval for using customer data in this specific model?
6. **Sensor Bias and Physical Responsibility Audit:** Have I verified that the system does not fail under changing environmental conditions (e.g. daylight vs. night) and do I know who is legally responsible for AI-driven physical actions?

Final reflection: Viability is not about whether you *can* do it, but whether you *should* do it.
An ethical AI is a trustworthy AI — and in business, trust is the hardest asset to regain once lost.

6. Conclusions

As a final conclusion to this exhaustive walkthrough of the five dimensions, it is clear that an AI project in 2026 is not a conventional IT project; it is a **cognitive re-engineering of the organization**. AI is the engine, but success lies in the design of the chassis (processes) and the quality of the fuel (data).

Here you have the strategic synthesis and global recommendations to move from planning to execution with guarantees.

C1. Final Reflection: The Role of AI as Copilot

The ultimate goal of preparing these five dimensions is **not** to replace the human, but to free them from low-value cognitive load. A well-prepared AI project turns your employees into "supervisors of intelligent systems", raising the quality bar of the entire company.

C2. The "Hybrid Orchestra" Paradigm

The most important technical conclusion is that no single tool does everything. The companies that succeed are those that build **hybrid architectures**:

- They use LLMs for interfaces and linguistic reasoning.
- They use classical Machine Learning for numerical predictions.
- They use agents and orchestration layers to connect systems.
- They keep humans in the critical decision nodes.

In practice, the organizations that win are the ones that treat AI like a hybrid orchestra: each model, system and human has a clear instrument to play, and the score —the process— is written with those constraints in mind.

C3. Roadmap: Global Recommendations

A. Start from the "Pain", not the "Hype"

Do not choose the most advanced technology (like Claude 4 or GPT-5) just because of its name. Identify the process step where human error is most expensive or the bottleneck is most suffocating. Economic viability begins by solving a real problem, not by creating an expensive experiment.

B. Build a "Data Culture" before a "Model Culture"

Invest more time in cleaning, labeling and connecting your data sources than in tweaking model parameters.
A mediocre model with excellent data will always outperform a brilliant model with noisy or biased data.

C. Design for the Exception

Complex automation fails at the edges. Your project will be only as strong as its ability to detect when it does not know what to do.

Define clear "route-to-human" protocols. User trust in AI is earned more by how it manages its errors than by its easy successes.

D. Incremental Implementation (The 80/20 Rule)

Do not try to automate 100% of the process on day one.
Aim to automate 80% of standard cases (the **Happy Path**).
This will give you immediate ROI and real data to train AI on the remaining 20% of complex cases.

C4. Final Launch Checklist (The "Go/No-Go")

Before pressing the start button, validate that the project meets the minimum thresholds in these six decision areas.
If it fails any of the critical No-Go criteria, the project should not be launched.

The table below summarizes, in compact form, the six decision areas.

Area	Key Questions	No-Go Criterion
Technical	Does the architecture support latency, volume, scalability?	The architecture cannot support the use case.
Operational	Is the process sufficiently defined, with fallback?	No clarity on steps, exceptions or recovery.
Data	Is there a reliable and traceable data foundation?	No minimally reliable source.
Integration & Action	Can AI act with safety and control?	Critical action has no human supervision.
Economic	Does ROI clearly outweigh total cost?	Error or maintenance cost invalidates the use case.
Ethics, Legal & Governance	Is the solution explainable and compliant?	No explainability, sponsor or accountability.

Below, each area with its minimum validations and No-Go conditions.

1. Technical

Question: Do I have an architecture that supports the latency, volume, data integrity and scalability the process requires?

Minimum validations:

- The model, agent or bot is correctly selected for the task.
- Latency is within the process's maximum acceptable time.
- The solution can scale without collapsing cost or performance.
- Stable integrations exist with required systems.
- There is observability, logging and recovery capability in case of failures.
- Cybersecurity coverage exists for server access, device access and model hacking.

Technical No-Go:

- The chosen architecture cannot support the projected volume.
- Latency makes the process unfeasible.
- There is no way to integrate the system with critical applications.
- Cybersecurity is not covered.

2. Operational

Question: Does the process map cover what to do when AI fails, drifts or encounters an unforeseen case?

Minimum validations:

- The process is broken down into clear steps.
- Edge cases and confidence thresholds are defined.
- A manual or semi-automatic fallback exists.
- There are clear owners for each phase.
- Error and recovery scenarios have been tested.
- There is a clear operating model: who monitors the system, who manages incidents, who reviews borderline cases and how often.

Operational No-Go:

- The process cannot be described step by step with sufficient clarity.
- There is no contingency plan if AI fails.
- Automation depends on everything going right all the time.
- There is no defined operating model (roles, schedules, incident owners).

3. Data

Question: Is there a minimum data foundation that is reliable, accessible, traceable and sufficient to operate?

Minimum validations:

- There is an identified source of truth for each critical data point.
- Data quality is acceptable: accuracy, completeness, timeliness and consistency.
- Required data are available in the right time and form.
- Required labeling or preprocessing is feasible.
- Data lineage and governance are understood.

Data No-Go:

- There is no minimally reliable source.
- Key data are fragmented or inconsistent across systems.
- The model would have to decide based on ambiguous information without clear business rules.
- Too much sensor noise prevents the model from reaching the minimum confidence threshold observed in field tests.

4. Integration and Action

Question: Can AI act safely, traceably and with the right level of autonomy?

Minimum validations:

- Tools, APIs or RPA are defined and constrained.
- Permissions follow the principle of least privilege.
- Autonomy limits are defined by type of action.
- There is state confirmation before closing a task.
- There is an auditable record of each action.

Integration No-Go:

- The final action has high impact and no human supervision.
- AI cannot confirm that the action was executed correctly.
- There is no control over what exactly the system can do.
- Mechanical risk is not mitigated with an independent hardware kill-switch.

5. Economic

Question: Does the expected value clearly exceed the total cost of building, operating and maintaining the solution?

Minimum validations:

- Cost per transaction is estimated.
- Expected ROI is quantified.
- Inference, maintenance, support and retraining costs are included.
- The use case has enough impact to justify the effort.
- There is budget to iterate and correct.

Economic No-Go:

- Error cost clearly exceeds the expected value of the use case.
- Operating cost grows faster than the benefit.
- The solution depends on an excessively expensive model without real need.

6. Ethics, Legal and Governance

Question: Can I explain how AI makes decisions and ensure it neither violates regulations nor generates unacceptable bias?

Minimum validations:

- The decision is explainable to the required level.
- There has been legal and compliance review.
- Biases and discrimination risks have been evaluated.
- Guardrails and shutdown mechanisms exist.
- A human owner of the solution has been defined.

Ethical/Legal No-Go:

- I cannot explain the decision logic in sensitive cases.
- The solution may bias outcomes without control.
- There is no business sponsor or operational owner.
- There is no clear accountability for AI decisions.

C5. What's Your Next Step?

The recommended next move is to start with a **reduced-scope pilot (MVP)** focused on the node of your process that has the cleanest data and the clearest success metrics. From there, you can iteratively expand coverage, complexity and autonomy as you gain confidence in the system.

7.Terms of Use and Contact

Intellectual Property and CC BY-NC 4.0 License

Collaboration and Projects

If your organization is exploring complex automation, generative AI solutions or the integration of intelligent agents in critical environments, I am open to collaborating on strategic projects.

We can work together from the early stages (use-case identification, risk analysis, architecture design) through to piloting and production rollout, always with a focus on operational resilience and governance.

- **LinkedIn:** www.linkedin.com/in/federico-jover
- **Email:** federico.jover.risk@gmail.com

Informal Contact

If you want to go deeper into these topics, I regularly have informal conversations to exchange experiences with people who are already in this space, are just getting started or are simply curious.

Get in touch with no obligation and we will surely have a good time discussing this world of AI. If you are an entrepreneur and just starting out, a mentoring conversation with no commitment might be helpful.